

Pattern Recognition and Machine Learning: Homework 11

Qingru Hu 2020012996

2023 年 5 月 16 日

1 Screenshots of each part of the code that you have finished

```
10 # Prepare the CIFAR-100 dataset and the dataloader
11 transform = transforms.Compose(
12     [transforms.ToTensor(),
13      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
14
15 batch_size = 128
16 trainset = torchvision.datasets.CIFAR100(root='./data', train=True,
17                                          download=True, transform=transform)
18 trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
19                                          shuffle=True, num_workers=4)
20
21 testset = torchvision.datasets.CIFAR100(root='./data', train=False,
22                                         download=True, transform=transform)
23 testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
24                                         shuffle=False, num_workers=4)
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43 class DownSamplingBlock(nn.Module):
44     def __init__(self, dim_in, dim):
45         super().__init__()
46         dim = int(2 * dim_in)
47         self.conv1 = nn.Conv2d(dim_in, dim, kernel_size=3, padding=1, stride=2)
48         self.bn1 = nn.BatchNorm2d(dim)
49         self.conv2 = nn.Conv2d(dim, dim, kernel_size=3, padding=1)
50         self.bn2 = nn.BatchNorm2d(dim)
51         self.shortcut = nn.Sequential(
52             nn.Conv2d(dim_in, dim, kernel_size=1, stride=2, bias=False),
53             nn.BatchNorm2d(dim)
54         )
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88 # Define the optimizer and loss function.
89 lossfunc = nn.CrossEntropyLoss()
90 optimizer = optim.Adam(net.parameters(), lr=0.001)
91
92 # Write the training loop.
93 for epoch in range(10):
94     epoch_loss = 0.0
95     for i, data in enumerate(trainloader, 0):
96         images, labels = data[0].cuda(), data[1].cuda()
97         optimizer.zero_grad()
98         outputs = net(images)
99         loss = lossfunc(outputs, labels)
100        loss.backward()
101        optimizer.step()
102        epoch_loss += loss.item()
103    # print epoch loss
104    print(f'Epoch: {epoch + 1}, Loss: {epoch_loss / 200:.3f}')
105
106 # Evaluate the model on the test dataset
107 correct = 0
108 total = 0
109 with torch.no_grad():
110     for data in testloader:
111         images, labels = data[0].cuda(), data[1].cuda()
112         outputs = net(images)
113         _, predicted = torch.max(outputs.data, 1)
114         total += labels.size(0)
115         correct += (predicted == labels).sum().item()
116 print(f'Accuracy of the ResNet on 10000 test images is {100*correct/total:.2f} %')
```

```
Files already downloaded and verified
Files already downloaded and verified
Epoch: 1, Loss: 7.528
Epoch: 2, Loss: 6.253
Epoch: 3, Loss: 5.459
Epoch: 4, Loss: 4.902
Epoch: 5, Loss: 4.488
Epoch: 6, Loss: 4.183
Epoch: 7, Loss: 3.927
Epoch: 8, Loss: 3.731
Epoch: 9, Loss: 3.544
Epoch: 10, Loss: 3.374
Accuracy of the ResNet on 10000 test images is 46.61 %.
```

2 The training loss and the test accuracy

3 The complete code of `main.py`

The complete code is in `main.py` under the same fold of the report.

Acknowledgement: Thank Yihan Zhou (周亦涵) for the discussion about this homework.